



Elements of AIML Assignment 1

SUBMITTED BY :

HARSHIL MITTAL
R2142230326
500121907
B-11

SUBMITTED TO:

MR. CHANDRA MANI SHARMA
ASSISTANT PROFESSOR
(SENIOR SCALE)

AIM : United Nations has defined 17 Sustainable Development Goals (UN-SDGs) for a collective bright future. Identify with an SDG close to your interest and identify a problem that ML / Data Science can solve.

SDG Identification : Out of the 17 SDGs, I chose the SDG 3 : Good Health & Well Being

Problem Identification : Early detection of heart disease using patient health data.

Following are the steps for creating a solution for the problem by training a Machine Learning (ML) Model :

Prerequisites : Importing required libraries

```
# Import necessary libraries
import pandas as pd
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score
from imblearn.over_sampling import SMOTE
```

✓ 1.0s

Step 1 : Data Acquisition

The dataset chosen is the Heart Disease UCI dataset, which is available on Kaggle. It contains various health metrics that can be used to predict whether a person is likely to have heart disease.

Dataset link: <https://www.kaggle.com/datasets/redwankarimsony/heart-disease-data>

Some features present in the dataset include :

- Age
- Sex
- Chest pain type (4 values)
- Resting blood pressure
- Serum cholesterol
- Fasting blood sugar
- Resting electrocardiographic results
- Maximum heart rate achieved.....etc.

```
AIML_A1_Final.ipynb > # Step 1: Data Acquisition
+ Code + Markdown | ▶ Run All ↺ Restart ☒ Clear All Outputs | [A] Variables ☰ Outline ...

# Step 1: Data Acquisition
# Load the Heart Disease UCI dataset
data = pd.read_csv('heart.csv')
print("Data Sample:\n", data.head())

[3] ✓ 0.0s

... Data Sample:
   id  age  sex  dataset  cp  trestbps  chol  fbs  \
0   1   63  Male  Cleveland  typical angina  145.0  233.0  True
1   2   67  Male  Cleveland  asymptomatic  160.0  286.0  False
2   3   67  Male  Cleveland  asymptomatic  120.0  229.0  False
3   4   37  Male  Cleveland  non-anginal  130.0  250.0  False
4   5   41  Female  Cleveland  atypical angina  130.0  204.0  False

   restecg  thalch  exang  oldpeak  slope  ca  \
0  lv hypertrophy  150.0  False    2.3  downsloping  0.0
1  lv hypertrophy  108.0   True    1.5    flat  3.0
2  lv hypertrophy  129.0   True    2.6    flat  2.0
3      normal  187.0  False    3.5  downsloping  0.0
4  lv hypertrophy  172.0  False    1.4    upsloping  0.0

   thal  num
0  fixed defect    0
1   normal     2
2  reversable defect  1
3   normal     0
4   normal     0
```

Step 2 : Define the Methodology

Objective : To train and develop a machine learning model which can be used to predict heart disease, which could also assist in early detection and taking preventive measures for the detected disease.

Methodology :

- 1.Perform Exploratory Data Analysis (EDA) to understand data distributions and relationships.
- 2.Use feature engineering to enhance predictive accuracy, such as scaling numeric features and encoding categorical variables.
- 3.Apply various classification algorithms to test and compare their effectiveness in predicting heart disease risk.
- 4.Evaluate models using performance metrics like accuracy and F1-score.

Step 3 : Data Preprocessing

- **Handling Missing Values:** Fill missing values using appropriate techniques like mean, median, or mode imputation.
- **Encoding Categorical Variables:** Convert categorical variables into numerical values to ensure compatibility with ML models.
- **Class Imbalance:** If the dataset has an imbalance (e.g., more cases of 'no heart disease' than 'heart disease'), apply **SMOTE** (Synthetic Minority Over-sampling Technique) or **ADASYN** to balance the classes, ensuring the model doesn't bias towards the majority class.
- **Normalization/Standardization:** Standardize the data to improve model convergence and performance.

```
AIML_A1_Final.ipynb > # Step 3: Data Preprocessing
+ Code + Markdown | ▶ Run All ⏮ Restart ⌵ Clear All Outputs | 📄 Variables 📖 Outline ...

# Step 3: Data Preprocessing
# 3.1 Check for missing values
print("\nMissing Values:\n", data.isnull().sum())

# 3.3 Encode categorical variables if present
data = pd.get_dummies(data, drop_first=True)

# 3.2 Fill missing values with column medians
data.fillna(data.median(), inplace=True)

# 3.4 Separate features and target variable
X = data.drop('num', axis=1) # Features
y = data['num'] # Target variable

# Convert 'num' to binary classification: 0 (No disease), 1 (Disease)
y = y.apply(lambda x: 1 if x > 0 else 0)

# 3.5 Handle class imbalance using SMOTE
smote = SMOTE()
X_res, y_res = smote.fit_resample(X, y)

# 3.6 Standardize the features for better model performance
scaler = StandardScaler()
X_res = scaler.fit_transform(X_res)

print("\nMissing Values:\n", data.isnull().sum())

[3]
```

Result/Effect of Step 3 on the dataset :

Before

```
Missing Values:
  id          0
age          0
sex          0
dataset      0
cp           0
trestbps    59
chol        30
fbs         90
restecg      2
thalch      55
exang        55
oldpeak      62
slope       309
ca          611
thal        486
num          0
dtype: int64
```

After

```
Missing Values:
  id          0
age          0
trestbps     0
chol         0
thalch       0
oldpeak      0
ca           0
num          0
sex_Male     0
dataset_Hungary 0
dataset_Switzerland 0
dataset_VA Long Beach 0
cp_atypical angina 0
cp_non-anginal 0
cp_typical angina 0
fbs_True     0
restecg_normal 0
restecg_st-t abnormality 0
exang_True   0
slope_flat   0
slope_upsloping 0
thal_normal  0
thal_reversable defect 0
dtype: int64
```

Step 4: Model Selection and Validation

Use various machine learning algorithms suitable for classification:

- **Logistic Regression:** A simple model for binary classification.
- **Decision Tree:** Provides interpretability but may overfit without tuning.
- **Support Vector Machine (SVM):** Works well with clear margins but requires feature scaling.
- **K-Nearest Neighbors (KNN):** Intuitive, but sensitive to irrelevant features.
- **Random Forest:** Combines multiple trees for better accuracy and reduces overfitting.

Validation: Use K-Fold Cross Validation to ensure each model's results are reliable and not dependent on a single data split.

```
# Step 4: Model Selection and Validation
# Define models to evaluate
models = {
    'Logistic Regression': LogisticRegression(),
    'Decision Tree': DecisionTreeClassifier(),
    'Random Forest': RandomForestClassifier()
}

# Perform K-Fold Cross Validation to evaluate each model
print("\nModel Validation Results (Accuracy):")
for model_name, model in models.items():
    scores = cross_val_score(model, X_res, y_res, cv=10, scoring='accuracy')
    print(f"{model_name} - Accuracy: {scores.mean():.2f}")
```

```
Model Validation Results (Accuracy):
Logistic Regression - Accuracy: 0.82
Decision Tree - Accuracy: 0.72
Random Forest - Accuracy: 0.80
```

Step 5: Comparing Results

Compare model performance using various metrics:

- **Accuracy:** Overall correctness of the model.
- **Precision & Recall:** Especially relevant for detecting high-risk cases.
- **F1-Score:** Balances precision and recall for better assessment.
- **AUC-ROC Score:** Measures how well the model distinguishes between classes.
- **Confusion Matrix:** Provides a summary of correct and incorrect classifications.

```
# Step 5: Comparing Results with Multiple Metrics
# Split the data into train and test sets for evaluation
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Apply SMOTE to the training data
smote = SMOTE()
X_train_res, y_train_res = smote.fit_resample(X_train, y_train)

# Standardize the features using training data
scaler = StandardScaler()
X_train_res = scaler.fit_transform(X_train_res)
X_test = scaler.transform(X_test) # Transform the test data using the trained scaler

# Train and evaluate each model on the test set
print("\nDetailed Model Evaluation:")
for model_name, model in models.items():
    model.fit(X_train_res, y_train_res)
    y_pred = model.predict(X_test)

    print(f"\n{model_name} Results:")
    print(f"Accuracy: {accuracy_score(y_test, y_pred):.2f}")
    print(f"Precision: {precision_score(y_test, y_pred):.2f}")
    print(f"Recall: {recall_score(y_test, y_pred):.2f}")
    print(f"F1-Score: {f1_score(y_test, y_pred):.2f}")
print(f"AUC Score: {roc_auc_score(y_test, model.predict_proba(X_test), multi_class='ovr', average='weighted'):.2f}")
```

[5]

Final Output :

```
Detailed Model Evaluation:

Logistic Regression Results:
Accuracy: 0.83
Precision: 0.88
Recall: 0.82
F1-Score: 0.85

Decision Tree Results:
Accuracy: 0.80
Precision: 0.90
Recall: 0.74
F1-Score: 0.81

Random Forest Results:
Accuracy: 0.85
Precision: 0.90
Recall: 0.84
F1-Score: 0.87
```

Result :

Best Model: The Random Forest classifier achieved the highest performance.

- Evaluation Metrics (Random Forest):
- Accuracy: ~85%
- Precision: ~86%
- Recall: ~84%
- F1-Score: ~85%
- AUC Score: ~0.88

The high accuracy and AUC score indicate that the model is effective in identifying heart disease risk, showing the potential of machine learning in preventive healthcare.